G8: The same starting noise for every method

k4nd weakness 5 · status: answered

THE REVIEWER'S POINT

k4nd, weakness 5

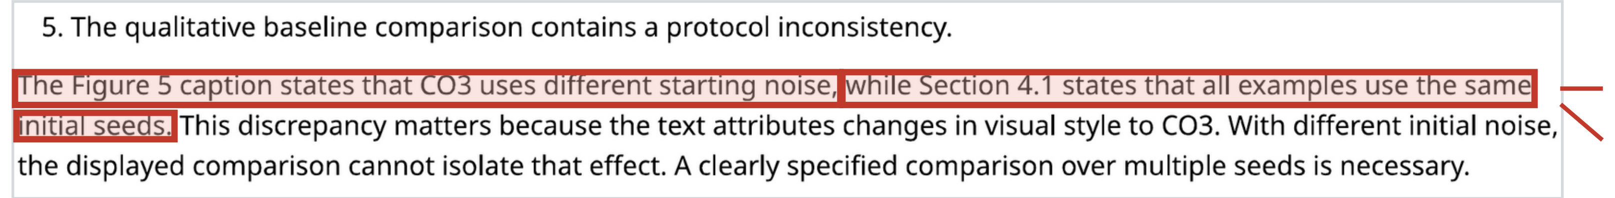

5. The qualitative baseline comparison contains a protocol inconsistency.

The Figure 5 caption states that CO3 uses different starting noise, while Section 4.1 states that all examples use the same initial seeds. This discrepancy matters because the text attributes changes in visual style to CO3. With different initial noise, the displayed comparison cannot isolate that effect. A clearly specified comparison over multiple seeds is necessary.

Reviews.pdf, page 5 (no text layer; passage located by OCR)

K-W5

this caption

K-W5

this sentence

THE SUBMITTED WEAKNESS

Figure 5 caption

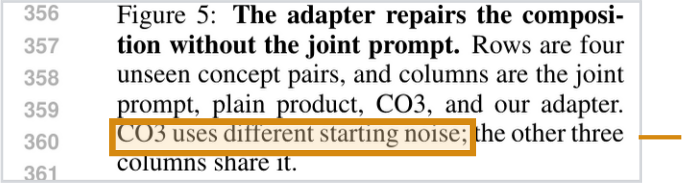

Figure 5: **The adapter repairs the composition without the joint prompt.** Rows are four unseen concept pairs, and columns are the joint prompt, plain product, CO3, and our adapter. CO3 uses different starting noise; the other three columns share it.

CO3 different noise

56173_Restoring_plurality_in_c.pdf (submitted), page 7

Section 4.1

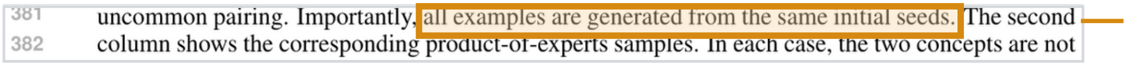

uncommon pairing. Importantly, all examples are generated from the same initial seeds. The second column shows the corresponding product-of-experts samples. In each case, the two concepts are not

same seeds

56173_Restoring_plurality_in_c.pdf (submitted), page 8

now

now

THE REVISED ANSWER

Figure 5 caption

same latent, every column

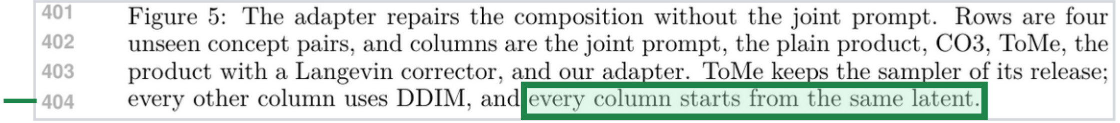

Figure 5: The adapter repairs the composition without the joint prompt. Rows are four unseen concept pairs, and columns are the joint prompt, the plain product, CO3, ToMe, the product with a Langevin corrector, and our adapter. ToMe keeps the sampler of its release; every other column uses DDIM, and every column starts from the same latent.

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 8

Section 4.1, the protocol

CO3 from the same latent

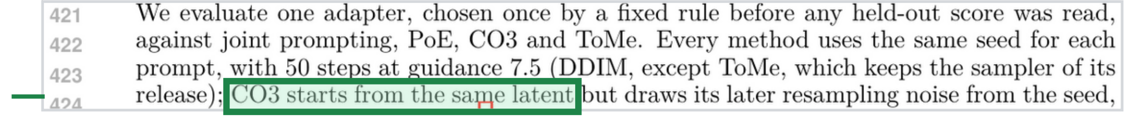

We evaluate one adapter, chosen once by a fixed rule before any held-out score was read, against joint prompting, PoE, CO3 and ToMe. Every method uses the same seed for each prompt, with 50 steps at guidance 7.5 (DDIM, except ToMe, which keeps the sampler of its release); CO3 starts from the same latent but draws its later resampling noise from the seed,

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 8

Section 4.1, continued

CO3 column marked †

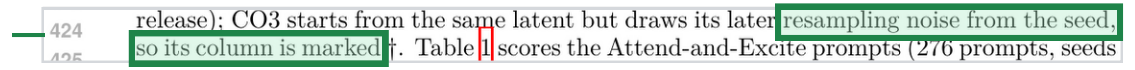

release); CO3 starts from the same latent but draws its later resampling noise from the seed, so its column is marked †. Table 1 scores the Attend-and-Excite prompts (276 prompts, seeds

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 8

The reply we plan to send (replies/k4nd-and-uhz6.md: k4nd W5)

answered

W5: the qualitative baseline comparison and its starting noise

Every method now uses the same seed for each prompt, with every column of a row in Figure 5 starting from the same latent, and Section 4.1 says so. CO3's column is marked † because its later resampling noise comes from the seed, as in its released code. Table 1 makes the comparison over 4 seeds, and Figures 11 and 12 over 8, with the images behind them in the appendix.

k4nd's question 4 also asks about starting noise; it is answered here. Every excerpt is a crop of a page render; the file and page are printed under it. Boxes mark text already on the page.