G4: Standard benchmarks and metrics

PrzE weakness 1 and question 1, UHz6 weakness 1 · status: answered

THE REVIEWER'S POINT

PrzE, weakness 1

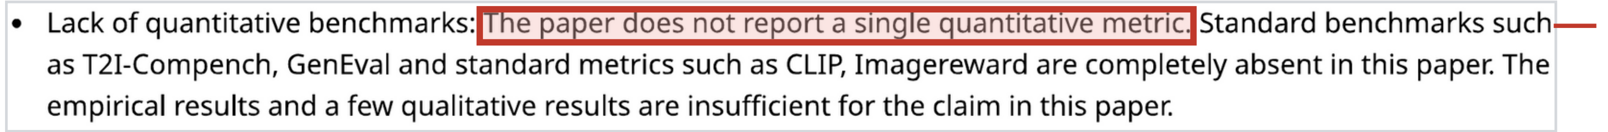

- Lack of quantitative benchmarks: The paper does not report a single quantitative metric. Standard benchmarks such as T2I-Compench, GenEval and standard metrics such as CLIP, Imagereward are completely absent in this paper. The empirical results and a few qualitative results are insufficient for the claim in this paper.

Reviews.pdf, page 2 (no text layer; passage located by OCR)

P-W1

no metric

PrzE, question 1

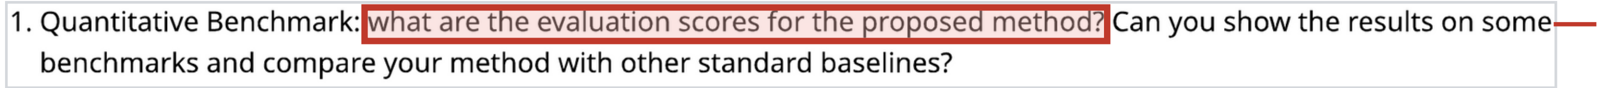

1. Quantitative Benchmark: what are the evaluation scores for the proposed method? Can you show the results on some benchmarks and compare your method with other standard baselines?

Reviews.pdf, page 2 (no text layer; passage located by OCR)

P-Q1

UHz6, weakness 1

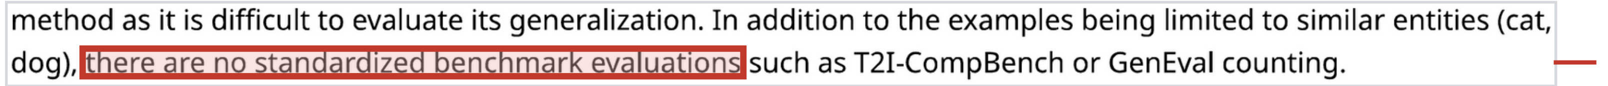

method as it is difficult to evaluate its generalization. In addition to the examples being limited to similar entities (cat, dog), there are no standardized benchmark evaluations such as T2I-CompBench or GenEval counting.

Reviews.pdf, page 3 (no text layer; passage located by OCR)

U-W1b

THE SUBMITTED WEAKNESS

Section 4.1, the results

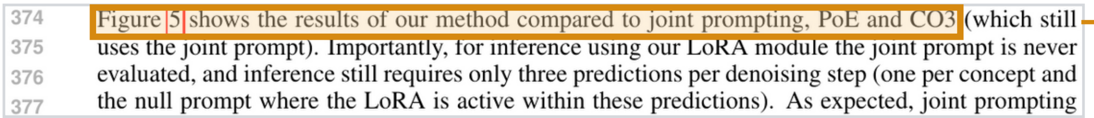

Figure 5 shows the results of our method compared to joint prompting, PoE and CO3 (which still uses the joint prompt). Importantly, for inference using our LoRA module the joint prompt is never evaluated, and inference still requires only three predictions per denoising step (one per concept and the null prompt where the LoRA is active within these predictions). As expected, joint prompting

results are image grids only

56173_Restoring_plurality_in_c.pdf (submitted), page 7

scored

benchmarks

THE REVISED ANSWER

Table 1, Attend-and-Excite prompts

CLIP, BLIP-VQA, ImageReward

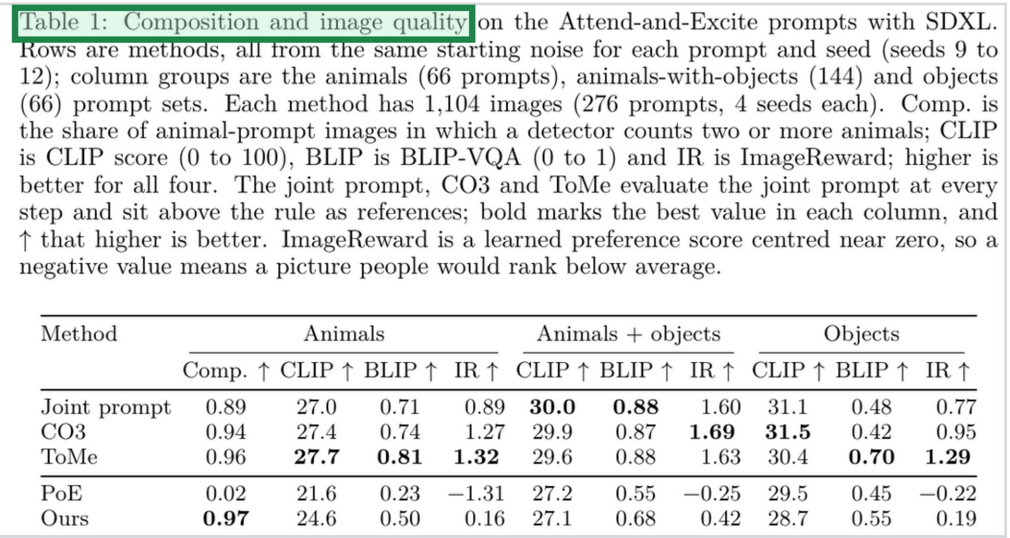

Table 1: Composition and image quality on the Attend-and-Excite prompts with SDXL. Rows are methods, all from the same starting noise for each prompt and seed (seeds 9 to 12); column groups are the animals (66 prompts), animals-with-objects (144) and objects (66) prompt sets. Each method has 1,104 images (276 prompts, 4 seeds each). Comp. is the share of animal-prompt images in which a detector counts two or more animals; CLIP is CLIP score (0 to 100), BLIP is BLIP-VQA (0 to 1) and IR is ImageReward; higher is better for all four. The joint prompt, CO3 and ToMe evaluate the joint prompt at every step and sit above the rule as references; bold marks the best value in each column, and ↑ that higher is better. ImageReward is a learned preference score centred near zero, so a negative value means a picture people would rank below average.

Method	Animals				Animals + objects			Objects		
	Comp. ↑	CLIP ↑	BLIP ↑	IR ↑	CLIP ↑	BLIP ↑	IR ↑	CLIP ↑	BLIP ↑	IR ↑
Joint prompt	0.89	27.0	0.71	0.89	**30.0**	**0.88**	1.60	31.1	0.48	0.77
CO3	0.94	27.4	0.74	1.27	29.9	0.87	**1.69**	**31.5**	0.42	0.95
ToMe	0.96	**27.7**	**0.81**	**1.32**	29.6	0.88	1.63	30.4	**0.70**	**1.29**
PoE	0.02	21.6	0.23	−1.31	27.2	0.55	−0.25	29.5	0.45	−0.22
Ours	**0.97**	24.6	0.50	0.16	27.1	0.68	0.42	28.7	0.55	0.19

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 9

Table 2, GenEval and T2I-CompBench

GenEval, T2I-CompBench

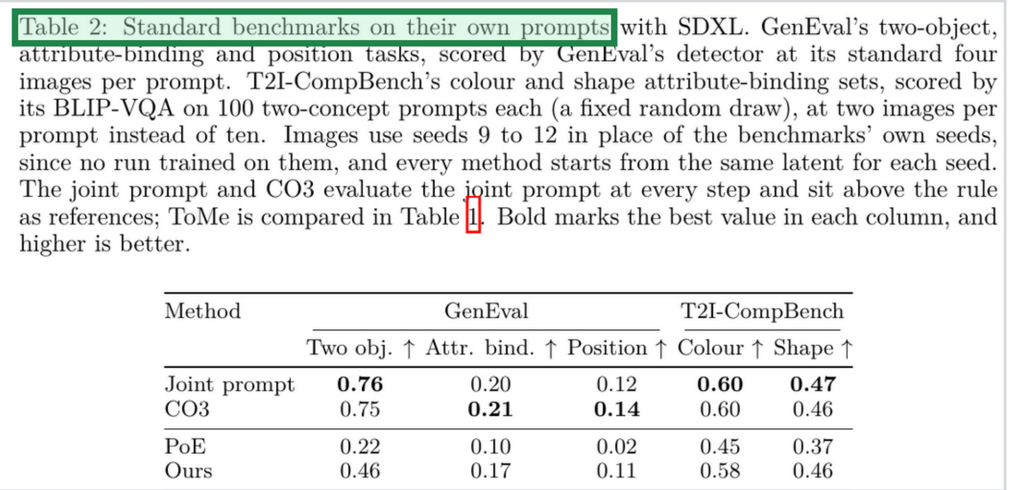

Table 2: Standard benchmarks on their own prompts with SDXL. GenEval's two-object, attribute-binding and position tasks, scored by GenEval's detector at its standard four images per prompt. T2I-CompBench's colour and shape attribute-binding sets, scored by its BLIP-VQA on 100 two-concept prompts each (a fixed random draw), at two images per prompt instead of ten. Images use seeds 9 to 12 in place of the benchmarks' own seeds, since no run trained on them, and every method starts from the same latent for each seed. The joint prompt and CO3 evaluate the joint prompt at every step and sit above the rule as references; ToMe is compared in Table 1. Bold marks the best value in each column, and higher is better.

Method	GenEval			T2I-CompBench	
	Two obj. ↑	Attr. bind. ↑	Position ↑	Colour ↑	Shape ↑
Joint prompt	**0.76**	0.20	0.12	**0.60**	**0.47**
CO3	0.75	**0.21**	**0.14**	0.60	0.46
PoE	0.22	0.10	0.02	0.45	0.37
Ours	0.46	0.17	0.11	0.58	0.46

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 15

The reply we plan to send (replies/prze.md: W1 and Q1)

answered

W1: lack of quantitative benchmarks

The revision adds quantitative evaluation. Table 1 scores every method on the 276 Attend-and-Excite prompts with BLIP-VQA (the T2I-CompBench composition metric), CLIP score and ImageReward. Table 2 reports GenEval's two-object, attribute-binding and position tasks and T2I-CompBench's colour and shape sets, scored by their official evaluators.

Q1: quantitative benchmark

Answered under W1 and W2: Tables 1 and 2.

Table 2 scores GenEval two-object, attribute binding and position, and T2I-CompBench colour and shape, each with its official evaluator.

The PrzE Q1 reply says "Tables 1 to 4"; the revision has three tables. UHz6's reply to weakness 1 is on the G3 page. Every excerpt is a crop of a page render; the file and page are printed under it. Boxes mark text already on the page.