G5: Baselines beyond CO3

PrzE weakness 2, UHz6 weakness 3 · status: partly answered (ToMe row of Table 1 pending)

THE REVIEWER'S POINT

PrzE, weakness 2

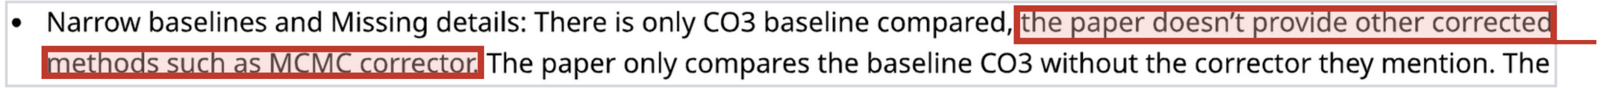

- Narrow baselines and Missing details: There is only CO3 baseline compared, the paper doesn't provide other corrected methods such as MCMC corrector. The paper only compares the baseline CO3 without the corrector they mention. The

Reviews.pdf, page 2 (no text layer; passage located by OCR)

P-W2a

only CO3

UHz6, weakness 3

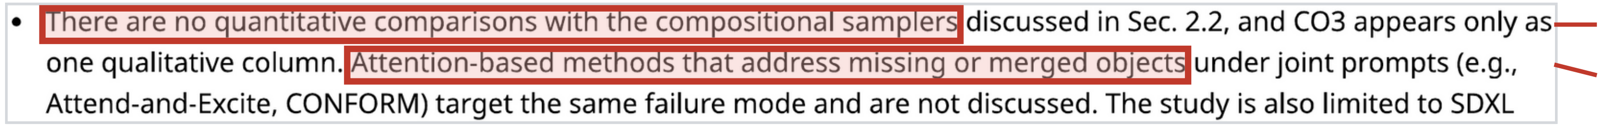

- There are no quantitative comparisons with the compositional samplers discussed in Sec. 2.2, and CO3 appears only as one qualitative column. Attention-based methods that address missing or merged objects under joint prompts (e.g., Attend-and-Excite, CONFORM) target the same failure mode and are not discussed. The study is also limited to SDXL

Reviews.pdf, page 3 (no text layer; passage located by OCR)

U-W3a

U-W3b

THE SUBMITTED WEAKNESS

Section 4.1, the results

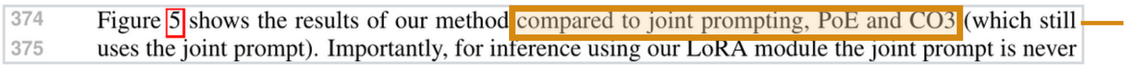

Figure 5 shows the results of our method compared to joint prompting, PoE and CO3 (which still uses the joint prompt). Importantly, for inference using our LoRA module the joint prompt is never

CO3 the only baseline

56173_Restoring_plurality_in_c.pdf (submitted), page 7

more columns

discussed

pending

THE REVISED ANSWER

Figure 5, six columns

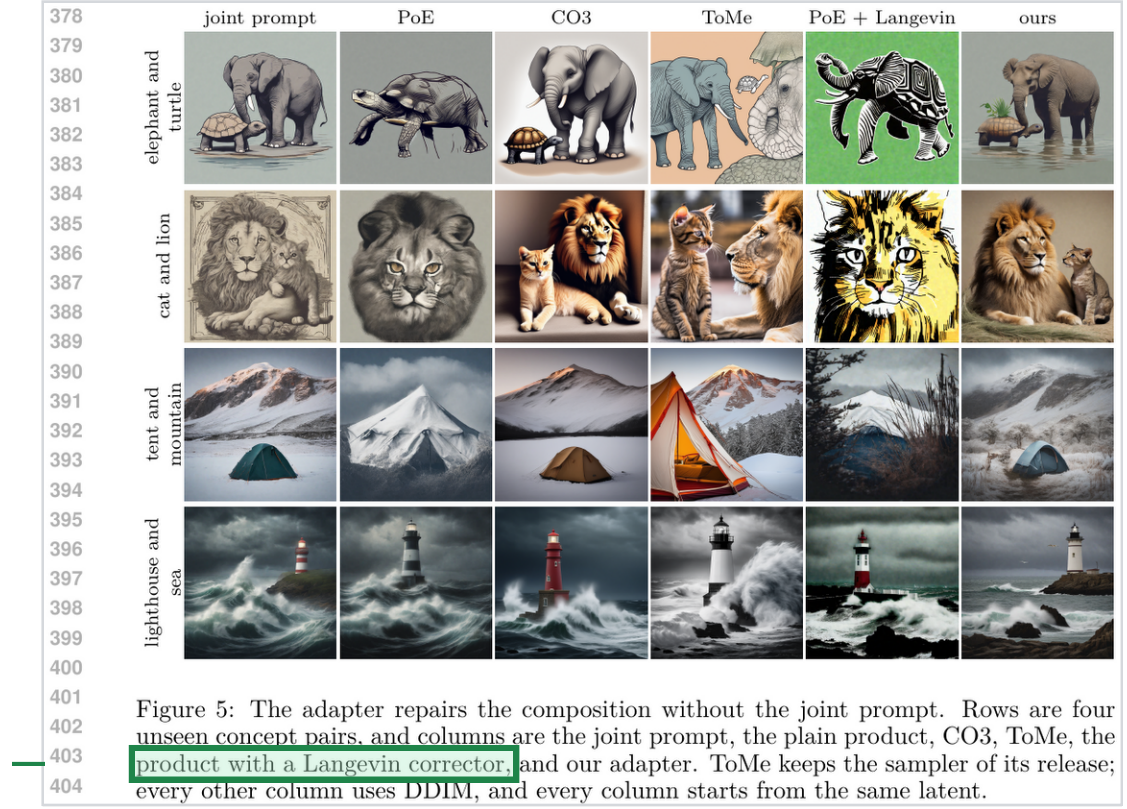

Figure 5: The adapter repairs the composition without the joint prompt. Rows are four unseen concept pairs, and columns are the joint prompt, the plain product, CO3, ToMe, the product with a Langevin corrector, and our adapter. ToMe keeps the sampler of its release; every other column uses DDIM, and every column starts from the same latent.

ToMe, Langevin added

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 8

Section 2, attention-based methods

now discussed

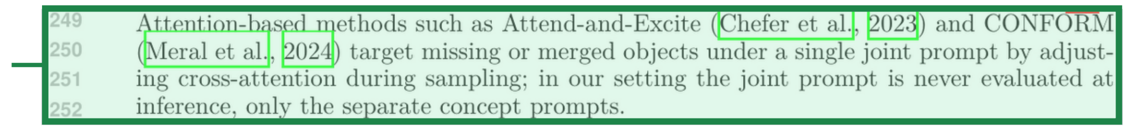

Attention-based methods such as Attend-and-Excite (Chefer et al., 2023) and CONFORM (Meral et al., 2024) target missing or merged objects under a single joint prompt by adjusting cross-attention during sampling; in our setting the joint prompt is never evaluated at inference, only the separate concept prompts.

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 5

Pending, and not in the revision

ToMe's row of Table 1 is rendering on the cluster (TODO in the PrzE reply).
The Langevin corrector appears in Figure 5 only, with no number in a table.

no crop: this panel states what the claim table records

The reply we plan to send (replies/prze.md: W2; replies/k4nd-and-uhz6.md: UHz6 W3)

partly answered

W2: narrow baselines and missing details

Figure 5 now includes a product of experts with a Langevin corrector (Du et al., 2023), the MCMC correction the paper discusses, run at its published setting from the same starting latent, beside CO3, ToMe and our adapter. Tables 1 and 3 compare with CO3 on every image. Table 1 adds ToMe, run on SDXL from the same starting latents [TODO: its row of Table 1]. Section 4 now describes the training set and how pairs were chosen (Figure 7); the 48 pairs are listed in our reply to Reviewer UHz6 and released with the code.

UHz6's weakness 3 reply (replies/k4nd-and-uhz6.md) says the same: CO3 in Tables 1 and 3, a Langevin corrector in Figure 5, Attend-and-Excite and CONFORM in Section 2.

Every excerpt is a crop of a page render; the file and page are printed under it. Boxes mark text already on the page.