G9: A missing figure, a broken sentence, repeated claims

UHz6 weakness 4, k4nd weakness 4 (last two parts) · status: answered

THE REVIEWER'S POINT

UHz6, weakness 4

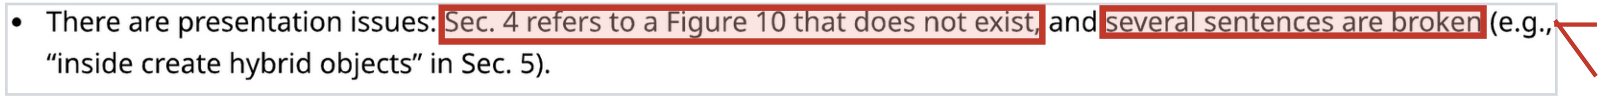

- There are presentation issues: Sec. 4 refers to a Figure 10 that does not exist, and several sentences are broken (e.g., "inside create hybrid objects" in Sec. 5).

Reviews.pdf, page 3 (no text layer; passage located by OCR)

U-W4a

U-W4b

k4nd, weakness 4 (continued)

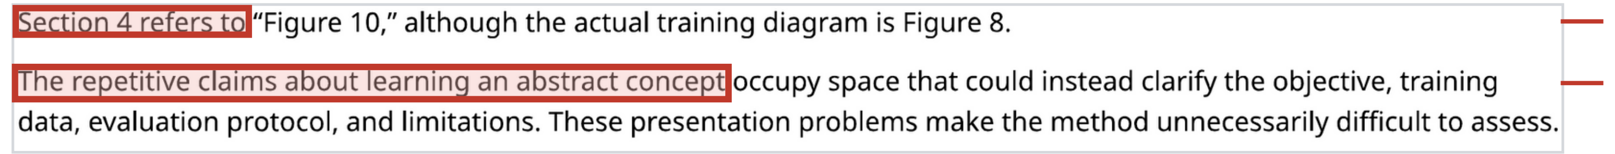

Section 4 refers to "Figure 10," although the actual training diagram is Figure 8.

The repetitive claims about learning an abstract concept occupy space that could instead clarify the objective, training data, evaluation protocol, and limitations. These presentation problems make the method unnecessarily difficult to assess.

Reviews.pdf, page 5 (no text layer; passage located by OCR)

K-W4c

K-W4d

this reference

repeated claim

this sentence

THE SUBMITTED WEAKNESS

Section 4, opening paragraph

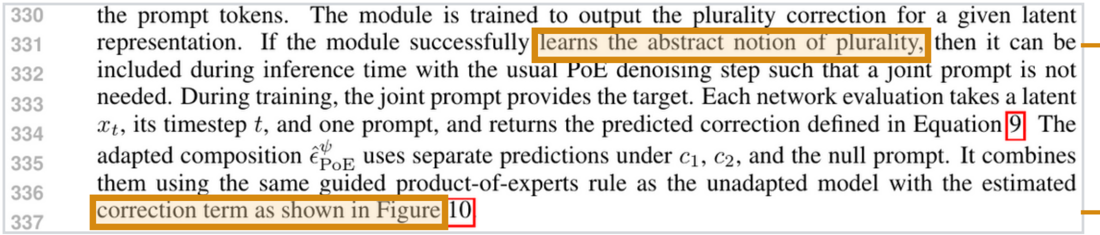

the prompt tokens. The module is trained to output the plurality correction for a given latent representation. If the module successfully learns the abstract notion of plurality, then it can be included during inference time with the usual PoE denoising step such that a joint prompt is not needed. During training, the joint prompt provides the target. Each network evaluation takes a latent x_t, its timestep t, and one prompt, and returns the predicted correction defined in Equation 9. The adapted composition $\hat{\epsilon}^{\psi}_{\mathrm{PoE}}$ uses separate predictions under c_1, c_2, and the null prompt. It combines them using the same guided product-of-experts rule as the unadapted model with the estimated correction term as shown in Figure 10

abstract notion of plurality

Figure 10 does not exist

56173_Restoring_plurality_in_c.pdf (submitted), page 7

Discussion

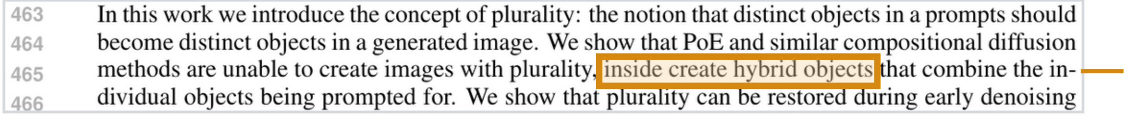

In this work we introduce the concept of plurality: the notion that distinct objects in a prompts should become distinct objects in a generated image. We show that PoE and similar compositional diffusion methods are unable to create images with plurality, inside create hybrid objects that combine the individual objects being prompted for. We show that plurality can be restored during early denoising

broken sentence

56173_Restoring_plurality_in_c.pdf (submitted), page 9

both removed

now reads

THE REVISED ANSWER

Section 4, opening paragraph

no Figure 10, no abstract notion

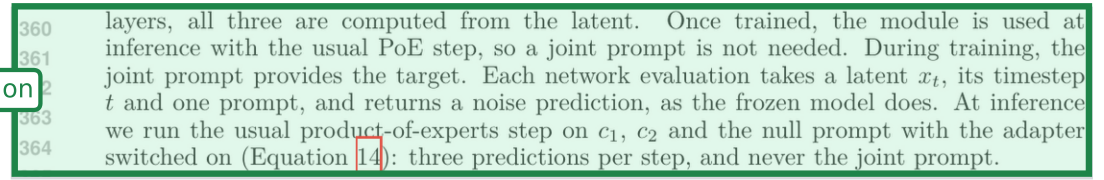

layers, all three are computed from the latent. Once trained, the module is used at inference with the usual PoE step, so a joint prompt is not needed. During training, the joint prompt provides the target. Each network evaluation takes a latent x_t, its timestep t and one prompt, and returns a noise prediction, as the frozen model does. At inference we run the usual product-of-experts step on c_1, c_2 and the null prompt with the adapter switched on (Equation 14): three predictions per step, and never the joint prompt.

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 7

Discussion

sentence repaired

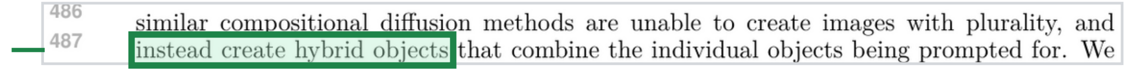

similar compositional diffusion methods are unable to create images with plurality, and instead create hybrid objects that combine the individual objects being prompted for. We

paper/iclr-56173/build/iclr2027_conference.pdf (revised), page 10

The reply we plan to send (replies/k4nd-and-uhz6.md: UHz6 presentation issues, k4nd W4 third part)

answered

Presentation issues

The reference that printed as "Figure 10" in Section 4 is removed, and the Discussion now reads "and instead create hybrid objects".

W4: method details poorly presented and sometimes inconsistent: Repetitive claims about learning an abstract concept

We removed the repeated claims about learning an abstract concept from Sections 3 and 4 and the Discussion. The introduction now states what the correction is and what it restores (see W1). Section 4 now describes the training data and how pairs were chosen (Figure 7). Appendix C.5 (Figure 15) and a new Discussion paragraph show the adapter on three concepts, including where it fails: it separates the concepts in every case but stops at two animals. Section 4.1 now opens with the evaluation protocol: one adapter chosen by a fixed rule, the same seed for every method, the prompt sets and held-out pairs with their seeds, the scorers, and the renders behind Table 1 in the appendix.

The repeated abstract-concept claims were cut from Sections 3 and 4 and the Discussion (06a9cf98); Section 4.1 opens with the evaluation protocol (508d2647). Every excerpt is a crop of a page render; the file and page are printed under it. Boxes mark text already on the page.